

MORTIS: EXPLOITING INTRA-PROMPT HETEROGENEITY FOR KV-CACHE MANAGEMENT IN LLM AGENT SERVING

Anonymous Authors¹

ABSTRACT

Existing agent serving systems schedule the KV-cache at the level of requests or prefixes. However, we observe that prompts in agent LLM invocations are not monolithic: their components differ in sharing scope, lifetime, and evolution pattern. Neglecting this heterogeneity leads to suboptimal KV-cache scheduling, but a system that sees only the token sequence can hardly identify and exploit it, because semantic boundaries and behavior differences are buried in the tokens. Our insight is that frameworks exposing model invocations through structured programming abstractions make this information available. We therefore present MORTIS, a compiler-serving co-design that analyzes the agent program to recover intra-prompt heterogeneity and exploits it for KV-cache scheduling. Our evaluation shows that MORTIS achieves up to $1.90\times$ lower average JCT, $4.49\times$ lower average TTFT, and $1.46\times$ higher throughput than Vanilla SGLang.

1 INTRODUCTION

LLM agents invoke a series of dependent model calls with shared context. Reusing the key-value (KV) cache across these calls avoids redundant prefill computation, making prompt reuse an important opportunity for efficient agent serving. Existing agent serving systems primarily optimize at the level of requests or LLM invocation sequences. Pythia (Yu et al., 2026) profiles agent workflows to guide cache management, request scheduling, and autoscaling. CacheScout (Zhang et al., 2026) learns execution transitions online to guide eviction and prefetching, while PBKV (Zheng et al., 2026) predicts future invocations and estimates the reuse potential of cached requests. These systems exploit predictability between invocations, but treat the contents of each prompt as an opaque token sequence.

However, we observe that prompts in agent LLM invocations are not monolithic: their components differ along three dimensions. (1) *Sharing scope*: some components appear in the prompts of many different agent steps, while others are specific to a single step. (2) *Lifetime*: some components stay valid for the whole session, while others become stale once a stage or loop iteration ends. (3) *Evolution pattern*: across invocations, a component may stay fixed, be replaced, or grow by appending. We call this **intra-prompt heterogeneity**.

Neglecting this intra-prompt heterogeneity can lead to sub-

optimal KV-cache management. Because mainstream serving engines reuse KV-cache by prefix matching, a change in an earlier component prevents the reuse of every token after it, even if those tokens are unchanged. Moreover, without knowing which components persist or evolve, the serving system cannot accurately determine which cached tokens should be retained, evicted, or prefetched. Figure 1 shows both problems in a four-stage financial-analysis agent: the stage-specific task instructions block the reuse of everything after them, and the serving engine cannot proactively release the KV-cache of the *documents* after Audit. However, it is hard for a system that sees only the token sequence to identify and exploit this intra-prompt heterogeneity, because semantic boundaries and reuse patterns are not represented in the tokens.




Figure 1. KV-cache reuse demonstration in a four-stage financial-analysis agent. This agent answers the question using the given documents. The stage-specific task instructions block KV-cache reuse, and the serving engine cannot proactively release the stale KV-cache (*documents* after Audit here) without knowing the intra-prompt heterogeneity.

Our key insight is that this information is available when

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

the prompt is programmatically generated. Agent frameworks increasingly expose LLM invocations through structured programming abstractions, such as byLLM (Dantananarayana et al., 2025), DSPy (Khatab et al., 2024), and NOAA (Furgale et al., 2026). These frameworks let developers declare each LLM invocation as a function with typed parameters, and the framework runtime generates the prompt from this declaration and the argument values. The sharing scope, lifetime, and evolution pattern of each component are therefore encoded in code, which can be systematically analyzed.

We therefore present MORTIS, a compiler-serving co-design that recovers intra-prompt heterogeneity from the agent program and exploits it for efficient KV-cache management. The compiler analyzes the program and passes the recovered heterogeneity to the serving system, which addresses both problems above: it determines prompt layouts that expose additional prefix reuse, and keeps the resulting KV states resident when they are likely to be needed. Identifying and exploiting intra-prompt heterogeneity introduces two challenges. **First, intra-prompt heterogeneity is buried in the interplay between program code and runtime values.** MORTIS recovers the sharing scope, lifetime, and evolution pattern of each prompt component through static analysis, and complements it with runtime observation of what the code cannot determine. **Second, KV-cache scheduling requires exposing opportunities under prefix caching and coordinating components with different heterogeneity characteristics.** To tackle this, MORTIS employs adaptive prompt re-layout to expose reusable prefixes, and a heterogeneity-aware scheduling strategy that guides KV-cache retention, eviction, and prefetching.

We implement MORTIS in Jac byLLM (Mars et al., 2023; Dantananarayana et al., 2025) and use SGLang (Zheng et al., 2024) as the underlying model-serving engine. We evaluate MORTIS on four representative agent workloads implemented in Jac byLLM: a HoVer-style (Jiang et al., 2020) fact-checking agent, a HumanEval (Chen et al., 2021) coding agent, a BFCL (Patil et al., 2025) tool-use agent, and a FinanceBench (Islam et al., 2023) financial-analysis agent. We further evaluate a mixed workload comprising all four agent types to assess MORTIS’s end-to-end performance under concurrent, heterogeneous demand.

In summary, we make the following contributions: (i) We identify the inherent heterogeneity in programmatically generated prompts and its implications for serving. (ii) We design and build MORTIS, a compiler-serving co-design that recovers and exploits prompt heterogeneity for efficient KV-cache management. (iii) We evaluate MORTIS against multiple established serving systems, demonstrating that MORTIS achieves up to $1.90\times$ lower average JCT and $4.49\times$ lower average TTFT than Vanilla SGLang, and

$1.46\times$ throughput on a mixed agent workload.

2 BACKGROUND

2.1 Agent frameworks with LLM calls as programming abstractions

A class of agent frameworks exposes model-backed computation through typed programming interfaces. Examples include byLLM’s meaning-typed programming (MTP) abstraction (Dantananarayana et al., 2025), DSPy’s Signatures and Modules (Khatab et al., 2024), and NOAA’s typed Python methods (Furgale et al., 2026). The developer declares a function or signature describing the desired behavior, inputs, and return type, and invokes it like an ordinary call within the program’s control flow. At each invocation, the framework runtime lowers the declaration and its argument values into model-facing messages.

2.2 KV-cache scheduling in LLM agent serving

LLM agent executions create an enormous KV-cache footprint, which should be carefully managed to avoid repeated computation. PagedAttention manages these states in shareable memory blocks (Kwon et al., 2023), while SGLang indexes cached token sequences in a radix tree (Zheng et al., 2024). An incoming request reuses the longest cached prefix of its prompt and computes KV for the remaining suffix.

Hierarchical caching systems such as CachedAttention (Gao et al., 2024), Mooncake (Qin et al., 2025), and SGLang Hi-Cache (SGLang Contributors, 2026) extend the KV-cache beyond device memory with host memory and, optionally, disk storage. KV evicted from the device can remain in the host tier and be reloaded on reuse, which can be cheaper than recomputation. However, reloading is not free. Transfers add latency unless they are overlapped with other work. Moreover, the host tier itself is finite. Offloading only relocates cache-capacity pressure but does not eliminate it.

Under memory pressure, engines evict KV at the granularity of whole requests or cached prefixes. SGLang evicts unused cached prefixes by recency (Zheng et al., 2024); Continuum (Li et al., 2025) pins KV across tool calls for a predicted duration; KVFlow (Pan et al., 2025) evicts first the KV farthest from execution, measured as steps-to-next invocation on the agent step graph; CacheScout (Zhang et al., 2026) prefetches and protects agent prompt anchors using transitions learned online; PBKV (Zheng et al., 2026) manages eviction and protection with predicted reuse from workflow history; Pythia (Yu et al., 2026) reconstructs the agent workflow from historical execution traces and uses the predicted execution graph for speculative KV-cache staging and eviction, request scheduling, and autoscaling across serving instances.

byLLM

```
system
This is a task you must complete by returning
only the output. ...

user
assess_evidence(search_results: str, question:
str, claim: str, evidences: list[Evidence]) ->
Finding --- (sem)
search_results: str ---- (sem)
question: str ---- (sem)
claim: str ---- (sem)
evidences: list[Evidence] ---- (sem)
search_results = (value)
question = (value)
claim = (value)
evidences = (value)
Schema requirements: (Finding fields)

response.format
(Finding JSON schema)
```

DSPy

```
system
Your input fields are:
1. 'search_results' (str): (desc)
2. 'question' (str): (desc)
3. 'claim' (str): (desc)
4. 'evidences' (list[Evidence]): (desc)
Your output fields are: 1. 'finding' (Finding):
(desc)
All interactions will be structured ...
[[ ## search_results ## ]] {search_results}
... [[ ## evidences ## ]] {evidences}
[[ ## finding ## ]] {finding} # note: ... JSON
schema: (Finding JSON schema)
[[ ## completed ## ]]
In adhering to this structure, your objective
is: (docstring)

user
[[ ## search_results ## ]] (value)
[[ ## question ## ]] (value)
[[ ## claim ## ]] (value)
[[ ## evidences ## ]] (value)
Respond with the corresponding output fields,
starting with '[[ ## finding ## ]]' ...
```

NVIDIA nooa

```
system
<strategy_prompt> You are generating structured
data matching the specified return type. ...
</strategy_prompt>

user
<sys tag="1"> Task(prompt='''# Your task
(docstring)
Args:
search_results: (arg doc)
question: (arg doc)
claim: (arg doc)
evidences: (arg doc)
## Input parameters:
search_results = (value)
question = (value)
claim = (value)
evidences = (value)
Perform the task and return the result
directly.'''> </sys>

response.format
(Finding JSON schema)
```

(sem), (docstring), (desc) and (arg doc) are the description strings each framework takes from the program (byLLM sem strings, DSPy Signature docstring and field desc, nooa function docstring and its Args entries); (value) is the runtime argument.

Figure 2. The prompt structure of programmatically generated prompts, lowered by three frameworks. They have similar structures: Same five kinds of component, only relative position differs.

3 MOTIVATION

Runtime-lowered prompts have similar structures Although these agent frameworks differ in surface syntax, they lower LLM invocations in similar ways. They declare LLM invocations as functions or function-like declarations, and the framework serializes that declaration together with the runtime arguments into a single prompt.

An example in Figure 2 demonstrates this. Jac byLLM, DSPy, and NOAA lowered prompts carry the same five kinds of components: framework scaffolding, task instruction, parameter descriptions, argument values, and output schema. The only difference is their relative position. Moreover, argument values are all emitted in declaration order.

The component order carries no intent The order in which these five components appear is decided solely by the framework runtime, without any program-specific knowledge. Figure 2 shows that the three representative frameworks arrange the same five components differently for the same invocation, which suggests that no particular order is required by the task. Even when argument values appear in declaration order, they do so simply because the runtime follows the parameter list, whose order was chosen by human developers.

Prompt components are heterogeneous A prompt programmatically generated by an agent framework is not prose. Instead, it has inherent heterogeneity. The runtime constructs the prompt from different sources, such as program constants, runtime values, and external inputs. Therefore

these components have different sharing scopes, lifetimes, and evolution patterns. As we discussed in §1 for Figure 1, this heterogeneity affects how the KV-cache can be managed effectively.

4 DESIGN

In this section, we present the design of MORTIS. MORTIS adopts a two-tier compiler-serving co-design. We define the *Agent Descriptor* as the interface to bridge the compiler and the server, which combines compiler analysis and runtime observation and is leveraged by the server side to enhance serving performance. As shown in Figure 3, MORTIS first identifies and annotates each component, and then efficiently plans the scheduling of the KV-cache.

4.1 Agent Descriptor

We define *Agent Descriptor* to bridge static analysis and server-side scheduling. The agent descriptor serves as a bridge between the two tiers of MORTIS, where static analysis provides a skeleton and the server refines the weights with runtime observations. An agent descriptor contains two parts: (i) **Call-site Template**: It provides a heuristic of the components’ lifetime, sharing scope and evolution pattern. (ii) **Workflow Automaton**: It contains the program-level structure of an agent, including workflow dependencies, execution time estimation and transition probabilities.

MORTIS traverses the agents type-checked AST to identify LLM call sites and their relationships. It then classifies each parameter’s provenance and lifetime to construct the




Figure 3. Architecture overview of MORTIS. Static analysis derives the agent descriptor, marking intra-prompt heterogeneity. The server uses it for KV-cache scheduling.

skeleton of the agent descriptor.

4.1.1 Call-site Template

As discussed in §3, the components of a programmatically generated prompt can be classified into the following two categories: (i) **Invariant Specifications**: the components that are fixed once the program is compiled, such as the system prompt, response schema, and the tool declarations. They are rendered from the compiler’s own intermediate representation and stay constant. (ii) **Dynamic Bindings**: because LLM calls are declared as ordinary functions, the arguments of a call site are supplied at runtime. These runtime values are the dynamic bindings.

We therefore define the call-site template as

$$\mathcal{T}_c = \langle \mathcal{I}_c, \mathcal{B}_c, \mathcal{H}_c \rangle,$$

where $\mathcal{I}_c$ is the invariant specifications, $\mathcal{B}_c$ is the dynamic bindings, and $\mathcal{H}_c$ is the heterogeneity annotation over those bindings. All three are recovered by static analysis.

Heterogeneity annotation. For each dynamic binding, $\mathcal{H}_c$ records both how its value evolves from one call to the next and how long that value stays valid:

$$\mathcal{H}_c : \mathcal{B}_c \longrightarrow \mathcal{E} \times \mathcal{S}.$$

The evolution component $\mathcal{E}$ classifies a binding by its value provenance. We distinguish $\mathcal{E}$ into five classes, listed from the most to the least reusable: (i) **CONST**: the value is fixed by the program text and is identical at every call. (ii) **COPY**: the value is the same as the value another call site carries, so the bytes are already resident once that call site has run. (iii) **RESP**: the value comes from the reply produced by another call site. (iv) **EXTEND**: the binding is an append-only context, and the previous value remains a prefix of the new one. (v) **VOLATILE**: the value is freshly produced at each call, or its origin lies outside what the analysis can see. We write ρ for the rank of a class in the order above, which expresses how much of a binding’s bytes an earlier call can already supply.

The lifetime component $\mathcal{S}$ names the program scope whose completion invalidates a binding’s value ($\mathcal{S} = \emptyset$ meaning that the value survives for the whole session once produced). Lifetime is independent of evolution, and the server needs both. The evolution class governs which bytes a later prompt can reuse, while the lifetime governs when cached bytes can be given up.

4.1.2 Workflow Automaton

An agent typically consists of a sequence of dependent LLM calls whose workflow structure is implicit in the program. Concretely, we abstract an agent as a time-annotated probabilistic finite-state machine (TPFM). We formulate the TPFM as:

$$\mathcal{A} = (Q, P, q_0, F, \mathcal{T}),$$

where: (i) Q is the set of states. Each state $q \in Q$ is labeled by an LLM callsite; (ii) P denotes the transition probabilities between states. (iii) $q_0 \in Q$ is the entrance callsite. (iv) F specifies the termination probability of the agent at each state. (v) $\mathcal{T} = \{D_q^{\text{exec}}, D_{q,q'}^{\text{gap}}\}$ denotes the timing annotations associated with states and transitions, where D_q^{exec} is the execution-time distribution of the LLM callsite represented by state q , and $D_{q,q'}^{\text{gap}}$ is the inter-call gap distribution between the completion of callsite q and the invocation of its successor q' .

Static analysis fixes the shape of this automaton, while its weights come from runtime logs. The states, the entry, the transitions among them, and final states are recovered from code. $\mathcal{T} = \{D_q^{\text{exec}}, D_{q,q'}^{\text{gap}}\}$ and P are often unknown at compile time and model-dependent. The server observes these online and accumulates them across sessions, updating the statistical model of the agent consistently. Together the two halves give the server a forward-looking view of a session.

4.2 Adaptive Prompt Re-layout

Prefix reuse stops at the first byte that differs from a cached prompt. The message a framework emits for a call site follows the declaration order of the function, which may place a changing value ahead of values that later calls could reuse. MORTIS’s adaptive re-layout permutes the components of each call message so that the bytes most likely to be resident already come first. The descriptor $\mathcal{T}_c$ (§4.1.1) and $\mathcal{A}$ (§4.1.2) determine a partial order over the components, and the server completes it into a layout with the token lengths it observes at runtime.

4.2.1 Movable Components

First, the dynamic bindings $\mathcal{B}_c$ can be permuted freely. The compiler lowers each argument into a self-contained component labeled by its parameter name, and the framework places these components in declaration order only because that is the order of the argument list; re-layout may therefore choose any order of the bindings.

Second, part of the invariant specification $\mathcal{I}_c$ is movable as well. $\mathcal{I}_c$ is a simplified abstraction. More precisely, it consists of the system prompt, the response schema, the tool schema, and the instruction specification (the task description and parameter descriptions generated by the compiler). Among them, the system prompt stays universally constant, and the two schemas are consumed by the model separately. The instruction specification is movable.

MORTIS therefore operates on the order of bindings and the placement of instruction specification. §4.2.2 describes how MORTIS chooses a layout within this space.

4.2.2 Layout Decision

Each call site sorts the bindings it carries by a program-wide key. Generally, *shared* bindings, which are carried by multiple call sites, are placed at the front, followed by the bindings carried by a single call site. Algorithm 1 summarizes the procedure.

Ordering private bindings. The bindings carried by a single call site are sorted in ascending lexicographic order of

$$K(b) = (\text{scoped}, \rho, \text{depth}, -\text{length}),$$

whose components are defined as follows: (i) *scoped*: the lifetime of b ; 0 if it is session-wide ($\mathcal{S} = \emptyset$) and 1 otherwise, so bindings that survive scope boundaries precede those reset by them. (ii) ρ : the stability rank defined in §4.1.1, from 1 for CONST to 5 for VOLATILE, so the more reusable class leads. (iii) *depth*: The number of transitions to the first call site that carries it from entrance, so values that enter a session earlier lead. (iv) *length*: the token length of the value of b , observed by the server; negated so longer

Algorithm 1 Layout decision algorithm

Require: Descriptor $(\{\mathcal{T}_c\}_{c \in C}, \mathcal{A})$, lengths ℓ , switch *pin*
Ensure: Layouts $\{L_c\}_{c \in C}$

- 1: $B_{\text{sh}} \leftarrow \{b : |\{c \in C : b \in \mathcal{B}_c\}| > 1\}; R \leftarrow \emptyset$
- 2: **for all** $e = (q, q') \in \text{EDGES}(\mathcal{A})$ **do**
- 3: $B_e \leftarrow B_{\text{sh}} \cap (\mathcal{B}_q \cup \mathcal{B}_{q'})$
- 4: $t_e(b) \leftarrow 1$ if *unchanged*, 2 if *append-only*, else 3
- 5: $R \leftarrow R \cup \{(b, b') \in B_e^2 : t_e(b) < t_e(b')\}$
- 6: $J(b) \leftarrow (-\ell(b), \text{depth}(b)); J_{\min}(S) \leftarrow \min_{b \in S} J(b)$
- 7: $D \leftarrow \text{CLUSTERCONFLICTS}(B_{\text{sh}}, R)$
- 8: $\langle S_1, \dots, S_k \rangle \leftarrow \text{TOPOSORT}(D, J_{\min}) \triangleright \text{Min-key first}$
- 9: $\pi_{\text{sh}} \leftarrow \text{SORT}(S_1, J) \parallel \dots \parallel \text{SORT}(S_k, J)$
- 10: **for all** $c \in C$ **do**
- 11: $\pi_c \leftarrow \langle b \in \pi_{\text{sh}} : b \in \mathcal{B}_c \rangle \parallel \text{SORT}(\mathcal{B}_c \setminus B_{\text{sh}}, K)$
- 12: $s_c \leftarrow \text{INSTRUCTIONSPEC}(\mathcal{I}_c)$
- 13: $L_c \leftarrow \langle s_c \rangle \parallel \pi_c$
- 14: **if** $\neg \text{pin} \wedge (\mathcal{B}_c \cap B_{\text{sh}} \neq \emptyset)$ **then** $L_c \leftarrow \pi_c \parallel \langle s_c \rangle$
- 15: **return** $\{L_c\}_{c \in C}$

bindings precede.

Ordering shared bindings. The ordering of shared bindings is derived from program-level structures. For every transition from q to q' , Static Analysis identifies how each shared binding of the two call sites evolves. We categorize these patterns into three groups, ordered as follows: (i) Both call sites carry it and its value does not change between the two calls. (ii) Both call sites carry it and the program only appends to its value. (iii) Program rewrites it between the two calls, or only one of the two call sites carries it.

MORTIS collects these requirements over all transitions and orders the shared bindings to satisfy them; when several bindings can go next, the one with the smallest (*-length, depth*) is placed first. Bindings with conflicting requirements are clustered and sorted by the same key.

Placing the specification. The *instruction specification* is lowered from code by the compiler and is unique to each call site. Emitting it first would therefore prevent prefix sharing among different call sites, however many values they have in common. After sorting, MORTIS moves the specification behind the bindings whenever the leading binding is shared, so that sibling call sites can share large blocks of common bindings. A per-program switch keeps the specification in front when moving it is suspected of changing the model’s behavior.

4.3 Proactive KV-Cache Scheduling

While adaptive prompt re-layout improves prefix reuse only by reshaping the request, MORTIS’s KV-cache scheduling proactively arranges the placement of KV-cache entries.

Using the agent descriptor, the server decides for every served KV-cache sequence what to keep, what to give up, and what to prefetch.

4.3.1 Reusable Head and Volatile Tail

Since re-layout already places the stable and shared bindings at the front, a served prompt naturally divides into two segments: a reusable head and a volatile tail. In a full prompt, everything behind the reusable head belongs to the volatile tail.

The reusable head consists of two parts: a *static prefix* and a *session-shared extension*: (i) **Static Prefix** is what the call site emits without runtime inputs, and its value outlives a single session. It includes the invariant specifications $\mathcal{I}_c$ together with the CONST bindings that re-layout has placed right behind them. (ii) **Session-shared Extension** is the run of bindings behind the static prefix that a later call of the same session will emit again at the same position. The two parts are derived from the agent program itself. $\mathcal{H}_c$ names the provenance of each binding, and the automaton $\mathcal{A}$ names the transitions between call sites. Together with the decided layout of each call site, we determine the reusable head by analyzing them jointly.

4.3.2 Next-Call Prediction

Because the workflow automaton $\mathcal{A}$ already exposes all potential next calls, our prediction focuses on the following three aspects: (i) The probability of the arrival. (ii) When the call will arrive. (iii) Which prefix of its prompt can be speculatively constructed.

Arrival probability. We estimate the probability of the next call by a loop-tagged table. As the compiler side emits the workflow graph, MORTIS adopts Bourdoncle-style (Bourdoncle, 1993) weak topological ordering to discover all loops within the agent, so that each call-site q is assigned the chain of loops enclosing it. The server counts the iterations of each loop per session, so every call carries a tuple ζ of the current iteration numbers of the loops enclosing its call-site, outermost first. We refine the transition and termination probabilities of $\mathcal{A}$ by this tuple and define the table as

$$\mathcal{P}(q' \mid q, \zeta) \quad \mathcal{F}(q, \zeta)$$

where ζ is the loop indices; $\mathcal{P}$ stands for the probability that a call-site q in context ζ is followed by a call-site q' ; $\mathcal{F}$ stands for the probability that the session ends at call-site q under context ζ . With an empty ζ they reduce to P and F of §4.1.2.

The table is a frequency count that the server updates across sessions. When a call arrives, its call-site is counted as the outcome of the previous call-site under the previous call's ζ , and termination is counted when the session closes.

Each observation is also counted under every suffix of ζ down to the empty tuple, so that for the same q the table holds one row per level of context, from the fully specified iteration to the pooled count over all iterations. After a call at q lands, we compute the probability of each outcome o by Witten–Bell interpolation from the empty context ζ_0 to the full context $\zeta_m = \zeta$, where o ranges over $\text{succ}(q)$, the static successors of q plus a virtual outcome end for session termination:

$$\hat{P}_k(o) = \lambda_k \frac{c_k(o)}{n_k} + (1 - \lambda_k) \hat{P}_{k-1}(o), \quad k = 0, \dots, m,$$

$$\hat{P}_{-1}(o) = \frac{1}{|\text{succ}(q)|}, \quad \lambda_k = \begin{cases} \frac{n_k}{n_k + d_k} & n_k > 0, \\ 0 & n_k = 0, \end{cases}$$

where ζ_k is the suffix of ζ that keeps only the innermost k indices, and the three counts are read from the row (q, ζ_k) of the table: $c_k(o)$ is how many times outcome o followed a call at q in context ζ_k , $n_k = \sum_o c_k(o)$ is the number of observations in that row, and d_k is the number of distinct outcomes with $c_k(o) > 0$.

Then,

$$\mathcal{P}(q' \mid q, \zeta) = \hat{P}_m(q') \quad \mathcal{F}(q, \zeta) = \hat{P}_m(\text{end})$$

A level with no observation is skipped. The weight λ_k grows with the support of a level, so a well-observed iteration dominates the prediction while a sparse one defers to the pooled behavior of q .

Arrival time. The time of arrival is conservatively approximated based on recorded statistics. The workflow automaton records $D_{q,q'}^{\text{gap}}$ per transition: each time a call at q' arrives, the server appends the interval between the reply of the preceding call at q and this arrival. The interval is the agent's own time spent in its code between two calls.

We estimate the gap before the next call as:

$$\hat{g} = t_{50} - \frac{1}{2}(t_{90} - t_{50})$$

where t_{50} and t_{90} are the 50th and 90th percentiles of the gap samples on the transition (q, q') ; the predicted call is then expected $\hat{g}$ after the reply of the current call. The second term ensures that a transition with a more dispersed gap distribution receives a tighter estimate.

Prefix reconstruction. As indicated in §4.3.1, the prompt of the known successor q' is reconstructible, or at least partially reconstructible. MORTIS builds its prompt in the layout $L_{q'}$, using the values that the session already carries. The construction processes bindings in the order of $L_{q'}$ while consulting their evolution class, deducing values accordingly. The construction stops at the first binding with no known origin. Then the matched prefix is sent to the kv-planner for speculative preparation.

4.3.3 Lifecycle Bands

As we discussed in §3, different components have different lifetimes, sharing scopes, and evolution patterns. Therefore, relying on recency-based LRU would lead to suboptimal performance. MORTIS adopts a banded priority and steers KV-cache eviction with it.

Bands. We categorize KV-cache entries into four bands, from first evicted to last: (1) **RETIRED**: the volatile tails of prompts belonging to completed sessions. (2) **TRANSIENT**: the volatile tail of a served prompt of a live session. The session that produced them will not present them again. (3) **UNCLASSIFIED**: entries that analysis cannot categorize, such as the prompts of a program that was never registered. They keep the engine’s default treatment. (4) **PLANNED**: prefixes that a predicted call will reuse. Inside this band entries are ordered by the expected value of keeping them resident, defined next.

Eviction empties the lowest band first and falls back to recency only within a band, so KV-cache that will likely be reused leaves memory later. On the host tier, however, it keeps its recency order except that retired segments go first, so a head the device gave up under pressure remains reloadable while it may still be needed. Additionally, KV-cache entries marked as *TRANSIENT* are not admitted into the host tier, as they would compete for memory with reusable KV-cache entries.

Band transitions. Bands are assigned at three moments of a session’s lifecycle. (1) **When a served prompt enters the cache**, MORTIS splits it at its reusable head (§4.3.1). The volatile tail and the reply generated after it become *TRANSIENT*; the head stays *UNCLASSIFIED* until the next prediction refresh finds a successor call that starts with it, which promotes it to *PLANNED*. (2) **When a later call arrives**, the transition it takes may rewrite bindings that earlier heads carried. Each affected head is cut back to the prefix that survives the rewrite, and the bytes past it are demoted to *TRANSIENT*. (3) **When the session ends**, the volatile tails of every prompt it served drop to *RETIRED*. Their heads stay *UNCLASSIFIED* and age out by recency, since other sessions of the same program may still share them.

These bands are mirrored directly into the radix tree’s eviction order. The band is the primary key of the eviction heap and recency breaks ties.

4.3.4 Retention and Prefetch

Based on the prediction, the kv-planner of MORTIS conducts protection. Its duty is to ensure that the required prefix resides on the device, or at least on the host. Each prediction is summarized into a tuple of reconstructed prefix, arrival

probability and estimated deadline:

$$\text{Prediction} = (\pi_{q'}, \mathcal{P}(q' | q, \zeta), t_{\text{reply}} + \hat{g})$$

where $\pi_{q'}$ is the prefix reconstructed for q' and t_{reply} is the time the current call replied. Predictions are refreshed at every call, and are dropped when the predicted calls arrive or expire. Protection acts on this tuple in two ways. *Retention* keeps $\pi_{q'}$ from leaving the device, and *prefetch* restores the part of $\pi_{q'}$ that has already been demoted to the host.

Retention. The planner ranks the predicted reusable KV-cache entries by the value of keeping $\pi_{q'}$ resident at every refresh:

$$\mathcal{D}(q') = \frac{\mathcal{P}(q' | q, \zeta) \cdot |\pi_{q'}|_{\text{cached}}}{\max(t_{\text{reply}} + \hat{g} - t_{\text{now}}, \tau_{\text{min}})}$$

where $|\pi_{q'}|_{\text{cached}}$ is the number of tokens of $\pi_{q'}$ cached (on either host or device) and τ_{min} is a floor that keeps a call due now from dominating. MORTIS’s kv-planner sorts all predictions by $\mathcal{D}$ in descending order, and admits them into *PLANNED* until the band covers a fixed fraction of the device tier. Within the band, their $\mathcal{D}$ add up on every radix cache node, so that a shared head outranks a private one.

Prefetch. When part of $\pi_{q'}$ has been demoted to the host, the planner prefetches it back to the device ahead of the estimated arrival time. A prefetch release deadline is therefore defined as:

$$t_{\text{release}} = t_{\text{reply}} + \hat{g} - \frac{|\pi_{q'}|_{\text{host}}}{B_{\text{h2d}}} - \frac{|\pi_{q'}|_{\text{uncached}}}{B_{\text{prefill}}}$$

where B_{h2d} and B_{prefill} are the measured host-to-device copy rate and prefill rate. Among released predictions, the earliest deadline goes first, with ties broken by expected saving:

$$\text{Saving} = \mathcal{P}(q' | q, \zeta) \cdot |\pi_{q'}|_{\text{cached}}$$

A prefetch takes only free device space after reclaiming *RETIRED* and *TRANSIENT* entries, so it never interferes with real requests; when no space is free, prefetching pauses briefly but still prevents predicted KV-cache entries from being completely evicted. A prefetch is released when the engine is idle or its deadline t_{release} approaches.

5 EVALUATION

In this section, we evaluate MORTIS on multiple representative agent workloads, and ablate each component.

5.1 Experimental Setup

5.1.1 Testbed and Model

We conduct all experiments on one NVIDIA RTX 3090 GPU using Qwen3-14B-AWQ. We allocate 5.2 GB of GPU

memory to the device KV-cache pool and 10 GB of CPU memory to the host KV-cache pool.

5.1.2 Baselines

We compare MORTIS against the following serving systems, together with two ablated variants of MORTIS. We also run each established serving system with MORTIS’s adaptive prompt re-layout enabled. All baselines are reimplemented on MORTIS’s SGLang fork with their original scheduling algorithms: (i) **Vanilla SGLang** (Zheng et al., 2024) uses SGLang’s HiCache with LRU eviction. (ii) **CacheScout** (Zhang et al., 2026) identifies agents from their request prefixes and learns inter-agent execution transitions online. It uses the learned transition patterns to prioritize cached KV entries and prefetch the prefixes of predicted successor agents. (iii) **Continuum** (Li et al., 2025) retains each agent program’s KV-cache for an adaptively selected time-to-live (TTL). The TTL is derived from observed inter-invocation delays and expected cache reuse, allowing the system to preserve reusable context while the agent executes tools. (iv) **KVFlow** (Pan et al., 2025) uses an Agent Step Graph to build a workflow-aware KV-cache management framework for agentic workloads. (v) **Static-Adapter** mirrors the adapter layer of DSPy (DSPy Contributors, 2026), which applies a fixed, program-agnostic layout to every call, hoisting the invariant components to the front. (vi) **MORTIS-KV-only** enables proactive KV-cache scheduling but disables adaptive prompt re-layout. (vii) **MORTIS-Relayout-only** enables adaptive prompt re-layout but replaces proactive KV-cache scheduling with LRU eviction.

5.1.3 Workloads

We evaluate MORTIS on four representative agent workloads, written in Jac byLLM (Dantanarayana et al., 2025): a HoVer-style (Jiang et al., 2020) fact-checking agent, a HumanEval (Chen et al., 2021) coding agent, a BFCL (Patil et al., 2025) tool-use agent, and a FinanceBench (Islam et al., 2023) financial-analysis agent. We further evaluate a mixed workload comprising all four agent types to assess MORTIS’s end-to-end performance under concurrent, heterogeneous demand. The mixed workload runs 16 concurrent lanes, four per agent type; when a lane finishes a session, it immediately starts the next one.

5.2 End-to-End Performance of MORTIS

We test MORTIS under different numbers of concurrent workflows and on a mixed workload. Across these workloads, we compare average job completion time (JCT) and average time-to-first-token (TTFT), and throughput (prompt tokens served per minute) for the mixed workload.

Table 1. Performance on the mixed workload. Metrics are measured over each system’s steady-state window.

System	Throughput	TTFT	Hit rate (%)		
	(tokens(k)/min)	(s)	Device	Host	Total
<i>Original prompt order</i>					
Vanilla SGLang	90.8	41.6	6.2	0.9	7.1
Static-Adapter	103.6	33.7	4.1	4.2	8.3
CacheScout	92.3	41.7	6.4	0.2	6.6
Continuum	90.6	42.1	4.4	1.0	5.4
KVFlow	92.6	41.2	7.5	0.6	8.1
MORTIS-KV-only	105.7	35.3	7.6	6.9	14.5
<i>With adaptive prompt re-layout</i>					
MORTIS-Relayout-only	108.8	35.6	5.0	7.5	12.6
CacheScout	113.2	34.4	9.4	5.3	14.7
Continuum	110.8	34.8	4.8	7.1	11.9
KVFlow	116.0	34.1	10.8	10.1	20.9
MORTIS w/o prefetch	129.3	29.7	9.3	20.9	30.2
MORTIS	132.3	29.0	9.1	22.6	31.6




Figure 4. Throughput and average TTFT speedup over Vanilla SGLang under varying load on the mixed workload.

Performance on mixed workload As shown in Table 1, MORTIS raises throughput to $1.46\times$ that of Vanilla SGLang and reduces average TTFT by $1.43\times$. Continuum and CacheScout perform close to Vanilla SGLang on all three metrics, and Static-Adapter stays within $1.14\times$ of it. The gain of MORTIS comes from a higher cache hit rate in both device and host memory.

We further evaluate MORTIS under varying pressure on the mixed workload. Keeping all four agents at the same lane count, we vary the number of lanes each agent holds. As shown in Figure 4, under low pressure adaptive prompt re-layout alone obtains nearly all of the gain, since device memory is still large enough to retain the reusable KV-cache. As pressure grows, the cache becomes contested and the gain of re-layout alone shrinks, while proactive KV-cache scheduling retains a clear advantage by protecting the reusable heads that LRU would evict. Under the heaviest pressure the two configurations converge toward the baseline, as the engine is completely saturated and there is no




Figure 5. Average TTFT speedup (top) and average JCT speedup (bottom) over Vanilla SGLang under varying numbers of concurrent workflows.

space to operate.

Per-workload sweeps Across the four workloads, MORTIS reduces average JCT by $1.06\text{--}1.90\times$ and average TTFT by $1.14\text{--}4.49\times$ compared with Vanilla SGLang (Figure 5). When concurrency is low, adaptive prompt re-layout accounts for nearly all of the gain. As concurrency grows, the gain of re-layout alone shrinks and proactive KV-cache scheduling carries it: it evicts volatile tails first and protects the reusable heads that predicted calls will need. On fact check, scheduling contributes from the lowest concurrency. On fact check and BFCL, and on finance at low concurrency, the TTFT reduction exceeds the JCT reduction, because a reused prefix skips its prefill, whereas JCT also includes decoding and tool execution, which the KV-cache does not shorten.

Except for a $1.09\text{--}1.22\times$ improvement on BFCL at $c = 4$, Continuum, CacheScout and KVFlow perform close to Vanilla SGLang on both metrics. They exploit predictability between invocations but treat each prompt as an opaque token sequence. In the original prompt order, volatile bindings often precede reusable ones, making prefixes hold little that later calls can reuse. MORTIS-KV-only, which schedules on the original prompt order, shows the same limit.

Task Quality We compare task accuracy between systems that serve the original prompts and systems that serve the re-laid prompts, using a paired permutation test over inputs (Table 2). Re-layout changes accuracy by at most 0.8 percentage points on any workload, with no consistent direction and no significant difference ($p \geq 0.66$). The performance

Table 2. Task accuracy (%) of systems serving the original and the re-laid prompts, and p -values of a paired permutation test.

Workload	Original	Re-laid	Δ (pp)	p
Fact check	20.0	20.4	+0.4	0.95
Coding	92.1	91.8	-0.2	0.66
BFCL	7.2	6.4	-0.8	1.00
Finance	34.6	35.4	+0.8	0.82

gains of MORTIS therefore do not come at the cost of task quality.

5.3 Cache Efficiency of MORTIS

Figure 6 shows the cache hit rate of prompt tokens at the same concurrency levels as Figure 5. On coding and BFCL, adaptive prompt re-layout alone brings the hit rate at low concurrency to that of MORTIS: it places the stable and shared bindings at the front, where they form a reusable head that later calls hit. As concurrency grows, the hit rate of MORTIS-Relayout-only drops quickly, because LRU eviction does not distinguish reusable heads from volatile tails. With proactive KV-cache scheduling, MORTIS keeps a higher hit rate than MORTIS-Relayout-only as concurrency grows, and on coding its hit rate stays nearly flat. On fact check, MORTIS’s hit rate exceeds that of MORTIS-Relayout-only from the lowest concurrency. Continuum and CacheScout stay close to Vanilla SGLang at every concurrency. MORTIS-KV-only stays well below MORTIS, since under the original prompt order the reusable head it can protect is short. The mixed workload shows the same pattern




Figure 6. Cache hit rate (%) of prompt tokens under varying numbers of concurrent workflows.

(Table 1).

5.4 Ablation Studies

We ablate MORTIS on the mixed workload by enabling its components one at a time (Table 1). MORTIS w/o prefetch keeps adaptive prompt re-layout and proactive KV-cache scheduling but disables prefetch. Figures 5 and 6 show MORTIS-KV-only and MORTIS-Relayout-only across concurrency levels.

Contribution of Each Mechanism Adaptive prompt re-layout alone raises throughput to $1.20\times$ that of Vanilla SGLang and reduces average TTFT by $1.17\times$. Static-Adapter raises throughput by $1.14\times$ and reduces average TTFT by $1.23\times$ on the mixed workload. On the per-workload sweeps, however, it matches MORTIS-Relayout-only on finance, Vanilla SGLang on fact check, and degrades on BFCL (Figure 6 and Figure 5), because the effect of program-agnostic prompt layout hoisting depends on the program itself.

Adding proactive KV-cache scheduling brings throughput to $1.42\times$ and the TTFT reduction to $1.40\times$, and lifts the hit rate from 12.6% to 30.2%. Prefetch adds little on this workload ($1.46\times$, $1.43\times$), since under this load the device pool is rarely free enough for a prefetch to take place.

Complementarity of Re-layout and Scheduling Adaptive prompt re-layout and proactive KV-cache scheduling depend on each other. Applied to the original prompt order, proactive KV-cache scheduling raises throughput by only $1.16\times$ (MORTIS-KV-only); applied after re-layout, it raises throughput by $1.22\times$ relative to MORTIS-Relayout-only. The combined $1.46\times$ gain is larger than the sum of the two individual gains. Re-layout places the stable and shared bindings at the front, so a served prompt divides into

a reusable head and a volatile tail, and scheduling gets opportunities to protect. Moreover, KV-cache scheduling and adaptive prompt re-layout are complementary. Applied on the re-laid prompts, CacheScout, Continuum and KVFlow all improve on MORTIS-Relayout-only (Table 1), and on the per-workload sweeps the gap widens with concurrency (Figure 5).

6 RELATED WORK

Frameworks for serving LLM agents. KVFlow (Pan et al., 2025) schedules KV by distance to the next invocation on an application-supplied agent step graph. Continuum (Li et al., 2025), CacheScout (Zhang et al., 2026), and PBKV (Zheng et al., 2026) learn retention signals from execution history, and Pythia (Yu et al., 2026) reconstructs the workflow from traces to drive cache staging, scheduling, and autoscaling. All of them manage whole requests or agent steps. MORTIS instead recovers the structure encoded in the program, and manages the KV-cache at intra-prompt granularity.

Prompt layout and reordering. Existing systems treat prompt layout as a fixed convention or delegate it to a model. SGLang’s frontend reorders prompt statements to lengthen the shared prefix, but relies on GPT-4 to choose the order, which its authors note is not yet reliable (Zheng et al., 2024). DSPy’s adapters emit components in one fixed order across all programs (DSPy Contributors, 2026), and NOAA argues for a three-region layout that places static context first (Furgale et al., 2026). These orders are chosen by convention or by an unsound rewrite. MORTIS derives the layout by systematically analyzing the program itself.

KV reuse beyond exact prefixes. Prompt Cache fixes module positions in a schema (Gim et al., 2024), and CacheBlend selectively recomputes tokens when concatenating cached chunks (Yao et al., 2025). Both relax positional exactness to gain reuse, at the cost of soundness or extra recomputation.

7 CONCLUSION

The exposed intra-prompt heterogeneity of programmatically generated prompts creates opportunities for optimization. MORTIS discovers these opportunities by static analysis assisted by runtime observations, and steers KV-cache scheduling with this information. Our evaluation demonstrates that MORTIS reduces average JCT by up to $1.90\times$ and average TTFT by up to $4.49\times$, and raises throughput by $1.46\times$ on a mixed workload.

REFERENCES

Bourdoncle, F. Efficient chaotic iteration strategies with widenings. In *Formal Methods in Programming and Their*

- Applications*, volume 735 of *Lecture Notes in Computer Science*, pp. 128–141. Springer, 1993. doi: 10.1007/BFb0039704.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. d. O., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., Ryder, N., Pavlov, M., Power, A., Kaiser, L., Bavarian, M., Winter, C., Tillet, P., Such, F. P., Cummings, D., Plappert, M., Chantzis, F., Barnes, E., Herbert-Voss, A., Guss, W. H., Nichol, A., Paino, A., Tezak, N., Tang, J., Babuschkin, I., Balaji, S., Jain, S., Saunders, W., Hesse, C., Carr, A. N., Leike, J., Achiam, J., Misra, V., Morikawa, E., Radford, A., Knight, M., Brundage, M., Murati, M., Mayer, K., Welinder, P., McGrew, B., Amodei, D., McCandlish, S., Sutskever, I., and Zaremba, W. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Dantnarayana, J. L., Kang, Y., Sivasothyathan, K., Clarke, C., Li, B., Kashmira, S., Flautner, K., Tang, L., and Mars, J. Mtp: A meaning-typed language abstraction for ai-integrated programming. *Proceedings of the ACM on Programming Languages*, 9(OOPSLA2):1176–1204, 2025. doi: 10.1145/3763092.
- DSPy Contributors. ChatAdapter: DSPy api reference. Online documentation, 2026. URL <https://dspy.ai/api/adapters/ChatAdapter/>. Accessed September 14, 2026.
- Furgale, P., Klingler, S., Nolan, J., Staats, M., Di Lorenzo, G., Martinez Abad, E., Schüller, C., Dinu, R., Devoto, A., Berard, P., Kaplun, G., Sarafian, E., Roveri, R., Derczynski, L., and Silveira Cabral, R. NVIDIA-labs OO Agents: Native Python Object-Oriented Agents, 2026. URL <https://arxiv.org/abs/2607.20709>.
- Gao, B., He, Z., Sharma, P., Kang, Q., Jevdjic, D., Deng, J., Yang, X., Yu, Z., and Zuo, P. Cost-efficient large language model serving for multi-turn conversations with CachedAttention. In *2024 USENIX Annual Technical Conference (USENIX ATC 24)*, pp. 111–126, Santa Clara, CA, 2024. USENIX Association. URL <https://www.usenix.org/conference/atc24/presentation/gao-b-in-cost>.
- Gim, I., Chen, G., seob Lee, S., Sarda, N., Khandelwal, A., and Zhong, L. Prompt Cache: Modular attention reuse for low-latency inference. In *Proceedings of Machine Learning and Systems*, volume 6, pp. 325–338, 2024. URL https://proceedings.mlsys.org/paper_files/paper/2024/hash/a66caa1703fe34705a4368c3014c1966-Abstract-Conference.html.
- Islam, P., Kannappan, A., Kiela, D., Qian, R., Scherrer, N., and Vidgen, B. Financebench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*, 2023.
- Jiang, Y., Bordia, S., Zhong, Z., Dognin, C., Singh, M., and Bansal, M. HoVer: A dataset for many-hop fact extraction and claim verification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 3441–3460, 2020.
- Khattab, O., Singhvi, A., Maheshwari, P., Zhang, Z., Santhanam, K., Vardhamanan, S., Haq, S., Sharma, A., Joshi, T. T., Moazam, H., Miller, H., Zaharia, M., and Potts, C. DSPy: Compiling Declarative Language Model Calls into State-of-the-Art Pipelines. In *International Conference on Learning Representations*, 2024.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626, New York, NY, USA, 2023. Association for Computing Machinery. doi: 10.1145/3600006.3613165. URL <https://arxiv.org/abs/2309.06180>.
- Li, H., He, R., Mang, Q., Zhang, Q., Mao, H., Chen, X., Zhou, H., Zhang, H., Cheung, A., Gonzalez, J., and Stoica, I. Continuum: Efficient and robust multi-turn llm agent scheduling with kv cache time-to-live, 2025.
- Mars, J., Kang, Y., Daynauth, R., Li, B., Mahendra, A., Flautner, K., and Tang, L. The jaseci programming paradigm and runtime stack: Building scale-out production applications easy and fast. *arXiv preprint arXiv:2305.09864*, 2023.
- Pan, Z., Patel, A., Hu, Z., Shen, Y., Guan, Y., Li, W.-L., Qin, L., Wang, Y., and Ding, Y. KVFlow: Efficient prefix caching for accelerating LLM-based multi-agent workflows, 2025. URL <https://arxiv.org/abs/2507.07400>.
- Patil, S. G., Mao, H., Yan, F., Ji, C. C.-J., Suresh, V., Stoica, I., and Gonzalez, J. E. The berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 48371–48392. PMLR, 2025.
- Qin, R., Li, Z., He, W., Cui, J., Ren, F., Zhang, M., Wu, Y., Zheng, W., and Xu, X. Mooncake: Trading more storage for less computation—a KVCache-centric architecture for serving LLM chatbot. In *23rd USENIX Conference*

on *File and Storage Technologies (FAST 25)*, pp. 155–170, Santa Clara, CA, 2025. USENIX Association. URL <https://www.usenix.org/conference/fast25/presentation/qin>.

SGLang Contributors. HiCache system design and optimization. SGLang documentation, 2026. URL https://docs.sglang.io/docs/advanced_features/hicache_design. Accessed September 22, 2026.

Yao, J., Li, H., Liu, Y., Ray, S., Cheng, Y., Zhang, Q., Du, K., Lu, S., and Jiang, J. CacheBlend: Fast large language model serving for RAG with cached knowledge fusion. In *Proceedings of the Twentieth European Conference on Computer Systems*, New York, NY, USA, 2025. Association for Computing Machinery. doi: 10.1145/3689031.3696098. URL <https://www.micros oft.com/en-us/research/uploads/prod/2024/09/eurosys25-final999.pdf>.

Yu, S., Shu, J., Ni, Y., Qian, K., Li, X., Wang, Y., Zhang, J., Xu, Z., Yang, S., Zhu, L., Zhai, E., Lu, Q., Xing, J., Lu, Y., Jin, X., Liu, X., and Xu, H. Pythia: Exploiting workflow predictability for efficient agent-native LLM serving, 2026. URL <https://arxiv.org/abs/2604.25899>.

Zhang, R., Kim, C., Feng, S., Du, K., Liu, Y., Zhong, Y., Ching, C.-W., Jiang, J., and Hu, L. Learning agent execution for kv-cache management in agentic serving, 2026.

Zheng, H., Fu, F., Wu, J., Yuan, B., Zhang, Y., Wang, H., Zhu, Y., Yan, X., and Jiang, J. Efficient serving for dynamic agent workflows with prediction-based KV-cache management, 2026. URL <https://arxiv.org/abs/2605.06472>.

Zheng, L., Yin, L., Xie, Z., Sun, C., Huang, J., Yu, C. H., Cao, S., Kozyrakis, C., Stoica, I., Gonzalez, J. E., Barrett, C., and Sheng, Y. Sglang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, volume 37, 2024.