

ANON

Anonymous Authors¹

ABSTRACT

LLM agent frameworks increasingly expose model invocations through structured programming abstractions, with framework runtimes lowering program-level constructs into concrete prompts. We observe that these programmatically generated prompts are inherently heterogeneous: their components exhibit different sharing scopes, lifetimes, and evolution patterns across invocations. However, existing serving systems treat prompts as opaque text, losing this structure and missing opportunities for more effective KV-cache reuse and management. To address the gap, we present ANON, a compiler-serving co-design that recovers and exploits this heterogeneity for efficient KV-cache management.

1 INTRODUCTION

LLM agents invoke a series of dependent model calls with share context. Reusing the key-value (KV) cache across these calls avoids redundant prefill computation, making prompt reuse an important opportunity for efficient agent serving. Existing agent serving systems primarily optimize at the level of requests or LLM invocation sequences. Pythia (Yu et al., 2026) profiles agent workflows to guide cache management, request scheduling, and autoscaling. CacheScout (Zhang et al., 2026) learns execution transitions online to guide eviction and prefetching, while PBKV (Zheng et al., 2026) predicts future invocations and estimates the reuse potential of cached requests. These systems exploit predictability between invocations, but treat the contents of each prompt as an opaque token sequence.

However, we observed that prompts in agent LLM invocations are not monolithic, they differ in three dimensions: (1) Sharing scope: one value can be shared across multiple invocations, or only within a single invocation. (2) Lifetime: Different bindings have different lifetimes. Some bindings are single-shot volatile inputs, while others may be reuse across the whole session. (3) Evolution pattern: context for LLM calls can evolve differently overtime, such as stay fixed, being replaced, or accumulating across invocations. We name this as **intra-prompt heterogeneity**.

Neglecting this intra-prompt heterogeneity can lead to sub-optimal KV-cache management. Consider a financial analysis agent that answers a question about a company’s filings

¹ Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the Machine Learning and Systems (MLSys) Conference. Do not distribute.




Figure 1. An example of intra-prompt heterogeneity in a financial analysis agent.

in four stages: Extract, Analyze, Audit, and Write (Figure 1). The prompt of each stage consists of a system prompt, a task description, the question, the documents, and the outputs of earlier stages. The system prompt is shared by every invocation, whereas the task description is replaced at every stage. The question and the documents remain fixed, but only the first three stages read the documents. The outputs of earlier stages accumulate as the session proceeds. Mainstream KV-cache serving engines are prefix-based: reuse stops at the first token that differs. Because the task description precedes everything else, these stages only share system prompt, and Analyze and Audit recompute the KV states of the documents although their contents remain unchanged. Moreover, without knowing which components persist or evolve, the serving system cannot accurately determine which cached tokens should be retained, evicted, or prefetched. In this example, it cannot tell that the documents can be released after Audit. Neither problem is easy to solve in the serving

system alone, because the semantic boundaries and reuse patterns are hard to identify from token sequences.

This information, however, is available when the prompt is programmatically generated. Agent frameworks increasingly expose LLM invocations through structured programming abstractions, such as byLLM (Dantamarayana et al., 2025), DSPy (Khatab et al., 2024), and NOOA (Furgale et al., 2026). These frameworks let developers declare each LLM invocation as a function with typed parameters, and the framework runtime generates the prompt from this declaration and the argument values. The sharing scope, lifetime, and evolution pattern of each component are therefore largely determined by the program, which can be systematically analyzed. This provides an opportunity to identify and exploit intra-prompt heterogeneity.

We present ANON, a compiler-serving co-design that recovers intra-prompt heterogeneity from the agent program and exploits it for efficient KV-cache management. The compiler analyzes the program and passes the recovered heterogeneity to the serving system, which addresses both problems above: it determines prompt layouts that expose additional prefix reuse, and keeps the resulting KV states resident when they are likely to be needed. Identifying and exploiting intra-prompt heterogeneity introduces two challenges:

First, heterogeneity is not always explicit in prompt. The sharing scope and lifetime of bindings are closely bonded with the program-level structure and actual runtime assignment. Determining which layout best serves prefix caching is therefore not trivial. ANON combines static analysis with runtime information to determine prompt layouts. The compiler derives a partial order over prompt components based on their sharing scopes and lifetimes. Then, ANON uses the actual token lengths of bindings observed from log to order components left unconstrained by this partial order, producing a concrete layout that improves prefix cache efficiency.

Second, exploiting future reuse requires coordinating cache retention and prefetching. Prompt components differ in when they will be reused and how much recomputation their cached prefixes can avoid. Retaining these prefixes consumes capacity needed by other requests, while prefetching too early occupies device memory and prefetching too late fails to hide transfer latency. Inappropriate KV-cache eviction would lead to unnecessary recomputation. ANON addressed this challenge by combining static analysis hints and runtime observation to instruct retention, eviction and prefetching.

We implement ANON in Jac ByLLM (Mars et al., 2023; Dantamarayana et al., 2025) and use SGLang (Zheng et al., 2024) as the underlying model-serving engine. We evaluate ANON on four representative agent workloads implemented

in ByLLM Jac (Mars et al., 2023): a HoVer-style (Jiang et al., 2020) fact-checking agent, a HumanEval (Chen et al., 2021) coding agent, a BFCL (Patil et al., 2025) tool-use agent, and a FinanceBench (Islam et al., 2023) financial-analysis agent. We further evaluate a mixed workload comprising all four agent types to assess ANON’s end-to-end performance under concurrent, heterogeneous demand.

We evaluate ANON on three representative agent workloads: a HoVer-style (Jiang et al., 2020) fact-checking agent, a HumanEval (Chen et al., 2021) coding agent, and a BFCL (Patil et al., 2025) tool-use agent. We further evaluate a mixed workload comprising all three agent types to assess ANON’s end-to-end performance under concurrent, heterogeneous demand.

In summary, we make the following contributions: (1) We identify the inherent heterogeneity in programmatically generated prompts and its implications for serving. (2) We design and build ANON, a compiler-serving co-design that recovers and exploits prompt heterogeneity for efficient KV-cache management. (3) We conduct comprehensive evaluations of ANON against multiple established serving systems, demonstrating that ANON achieves up to 48% lower average JCT and 78% lower average TTFT than Vanilla SGLang, and $1.44\times$ throughput on a mixed agent workload.

2 BACKGROUND

2.1 Serving systems for LLM agents

LLM serving engines retain the attention keys and values (KV) of tokens they have already processed to avoid repeating computation. PagedAttention manages these states in shareable memory blocks (Kwon et al., 2023), while SGLang indexes cached token sequences in a radix tree (Zheng et al., 2024). An incoming request reuses the longest cached prefix of its prompt and computes KV for the remaining suffix. Reuse stops at the first token mismatch: that token and every subsequent token must be recomputed, even if identical text appears later in a cached sequence, because its preceding context has changed. Prompt Cache supports modular reuse beyond a shared prefix, but requires modules to retain fixed, schema-assigned token positions (Gim et al., 2024).

Retained KV competes with running requests for device memory. Under memory pressure, engines such as SGLang evict unused cached prefixes by recency (Zheng et al., 2024). Hierarchical caching systems, including CachedAttention (Gao et al., 2024), Mooncake (Qin et al., 2025), and SGLang HiCache (SGLang Contributors, 2026), extend the cache beyond device memory with host memory and, optionally, storage. KV evicted from the device can remain in the host tier and be transferred back when needed; prefetching can overlap this transfer with other work. Reloading




Figure 2. The prompt structure of LLM invocation function `assess_evidence`, lowered by three frameworks. They have similar structures: Same five kinds of component, only relative position differs.

retained KV can be cheaper than recomputing it. The host tier is also limited: admitting one sequence consumes space that could retain another, so offloading does not eliminate cache-capacity pressure.

An agent *session* is one execution of an agent task, comprising a series of dependent LLM calls separated by tool and code execution. Between calls, the session’s KV sits idle and is exposed to recency eviction precisely while the agent is doing work that may lead to its reuse. Agent-aware serving systems use execution information to manage these gaps. KVFlow uses an agent step graph to estimate steps to future invocations and guide eviction and prefetching (Pan et al., 2025). Continuum retains KV across tool calls with a time-to-live (Li et al., 2025). CacheScout learns transitions between agents online to guide retention and prefetching (Zhang et al., 2026), while PBKV predicts future agent invocations from workflow history and the current context (Zheng et al., 2026). These systems share a request- or agent-step-level view of future execution: they organize cache management around the reuse of whole requests or agent steps.

2.2 Agent frameworks with LLM calls as programming abstractions

A class of agent frameworks exposes model-backed computation through typed programming interfaces. Examples include byLLM’s meaning-typed programming (MTP) abstraction (Dantamarayana et al., 2025), DSPy’s Signatures and Modules (Khatab et al., 2024), and NOAA’s typed

Python methods (Furgale et al., 2026). SIGIL likewise compiles skill descriptions into harnesses containing typed model operations (Dantamarayana et al., 2026). The developer declares a function or signature describing the desired behavior, its inputs, and its return type; invoking it is an ordinary call within the program’s control flow. We use *call site* to refer these function invocation locations.

At each invocation, the framework runtime lowers the declaration and its argument values into model-facing messages. The declaration supplies the function name, docstring or semantic strings, parameter types and descriptions, and return schema. This information is known when the program is compiled: byLLM captures it in MTP’s meaning-based intermediate representation, MT-IR (Dantamarayana et al., 2025). The argument values become available only at runtime, when the call executes. Thus, a generated prompt combines a description fixed by the program with values supplied by that execution.

Some frameworks make deliberate choices about how to arrange this information. NOAA places static context first, appended event history next, and volatile state last to maximize KV-cache reuse across turns (Furgale et al., 2026). DSPy’s adapters use a fixed ordering of system instructions, demonstrations, conversation history, and current inputs; its ChatAdapter emits input fields in signature order (DSPy Contributors, 2026). These layouts follow hand-chosen conventions applied across call sites and programs, rather than an order derived from the properties of a particular program’s bindings. The framework knows the structure

behind every prompt, but the serving engine receives only the token sequence.

3 MOTIVATION

Runtime-lowered prompts have similar structures Although these Program-Integrated LLM frameworks differ in surface syntax, they lower LLM invocations in similar ways. They declare LLM invocations as functions or function-like declarations, and framework serialize that declaration together with the runtime arguments into a single prompt.

An example in Figure 2 demonstrate this. ByLLM in Jac, DSPy, NOAA lower the same `assess_evidence` LLM invocation. These prompts carry the same five kinds of components: framework scaffolding, task instruction, parameter descriptions, argument values, and output schema. The only difference is their relative position. What’s more, argument values are all emitted in declaration order.

The component order carries no intent The order in which these five components appear is decided solely by the framework runtime, without any program specific knowledge. Figure 2 showed the three representative frameworks arrange the same five components differently for the same equivalent invocation, which suggests that no particular order is required by the task. Even when argument values appear in declaration order, they do so simply because the runtime follows the parameter list, whose order was chosen by human developers.

Prompt Components are heterogeneous A prompt programmatically generated by an agent framework is not prose. Instead, it has inherent heterogeneity. The runtime constructs the prompt from different sources, such as program constants, runtime values, external inputs. Therefore these components have different sharing scopes, lifetimes and evolution pattern. As we discussed in section 1 for Figure 1, this heterogeneity affects how the KV-cache can be managed effectively.

In summary, the heterogeneity within these programmatically generated prompt exposes new opportunities for efficient serving. However, none of these informations are served for effective serving, which motivates the design of ANON: Recovers heterogeneity by static analysis and use it to unlock more efficient KV-cache scheduling.

4 DESIGN

4.1 Overview

In this section, We present the design of ANON, an efficient serving engine for program-integrated LLM agents. ANON adopts a two-tier compiler-server co-design. On compiler

side, ANON performs static analysis to gather heterogeneity metadata, and on the server side, ANON leverages this metadata to optimize KV-cache scheduling. We first introduce an interface *Agent Descriptor* to bridge the compiler and server, together with the methodology of their construction. Second, we present a heterogeneity-aware KV-cache scheduling strategy that leverages the agent descriptor to enhance serving performance.

4.2 Agent Descriptor for Heterogeneity Metadata

Here we define *Agent Descriptor* to bridge static analysis and server-side scheduling. An agent descriptor contains two parts: (i) **Call-site Template**: It provides a heuristic of the components’ lifetime, sharing scope and evolution pattern. (ii) **Workflow Automaton**: It contains the program-level structure of an agent, including workflow dependencies, execution time estimation and transition probabilities. ANON’s two-tier design let the compiler to provide a skeleton, while the server refines the weight with runtime observations. Two sides are therefore connected by this agent descriptor, and jointly serve a more efficient KV-cache scheduling strategy.

4.2.1 Call-site Template

As discussed in §3, the components of a programmatically generated prompt can be classified into the following two categories: (i) **Invariant Specifications**: the components that are fixed once the program is compiled, such as the system prompt, response schema, and the tool declarations. They are rendered from the compiler’s own intermediate representation and stays constant. (ii) **Dynamic Bindings**: because LLM calls are declared as ordinary functions, the arguments of a call site are supplied in runtime. These runtime values are the dynamic bindings.

We therefore define the call-site template as

$$\mathcal{T}_c = \langle \mathcal{I}_c, \mathcal{B}_c, \mathcal{H}_c \rangle,$$

where $\mathcal{I}_c$ is the invariant specifications, $\mathcal{B}_c$ is the dynamic bindings, and $\mathcal{H}_c$ is the heterogeneity annotation over those bindings. All three are recovered by static analysis.

Heterogeneity annotation. For each dynamic binding, $\mathcal{H}_c$ records both how its value evolves from one call to the next and how long that value stays valid:

$$\mathcal{H}_c : \mathcal{B}_c \longrightarrow \mathcal{E} \times \mathcal{S}.$$

The evolution component $\mathcal{E}$ classifies a binding by its value provenance. We distinguish $\mathcal{E}$ into five classes, listed from the most to the least reusable: (i) **CONST**: the value is fixed by the program text and is identical at every call. (ii) **COPY**: the value is the same as the value another call site carries,




Figure 3. Architure overview of ANON.

so the bytes are already resident once that call site has run. (iii) **RESP**: the value comes from the reply produced by another call site. (iv) **EXTEND**: the binding is an append-only context, and the previous value remains a prefix of the new one. (v) **VOLATILE**: the value is freshly produced at each call, or its origin lies outside what the analysis can see. We write ρ for the rank of a class in the order above, which expresses how much of a binding’s bytes an earlier call can already supply.

The lifetime component $\mathcal{S}$ names the program scope whose completion invalidates a binding’s value ($\mathcal{S} = \emptyset$ meaning that the value survives for the whole session once produced). Lifetime is independent of evolution, and the server needs both. The evolution class governs which bytes a later prompt can reuse, while the lifetime governs when cached bytes can be given up.

4.2.2 Workflow Automaton

An agent typically consists of a sequence of dependent LLM calls whose workflow structure is implicit in the program. Concretely, we abstract an agent as a time-annotated probabilistic finite-state machine (TPFM). We formulate the TPFM as:

$$\mathcal{A} = (Q, P, q_0, F, \mathcal{T}),$$

where:

1. Q is the set of states. Each state $q \in Q$ is labelled by an LLM callsite;
2. P denotes the transition probabilities between states.

3. $q_0 \in Q$ is the entrance callsite.

4. F specifies the termination probability of the agent at each state.

5. $\mathcal{T} = \{D_q^{\text{exec}}, D_{q,q'}^{\text{gap}}\}$ denotes the timing annotations associated with states and transitions, where D_q^{exec} is the execution-time distribution of the LLM callsite represented by state q , and $D_{q,q'}^{\text{gap}}$ is the inter-call gap distribution between the completion of callsite q and the invocation of its successor q' .

Static analysis fixes the shape of this automaton, while its weights comes from runtime logs. The states, the entry, the transitions among them, and final states are recovered from code. $\mathcal{T} = \{D_q^{\text{exec}}, D_{q,q'}^{\text{gap}}\}$ and P are often unknown at compile time and model dependent. The server observes these online and accumulates them across sessions, updating the statistical model of the agent consistently. Together the two halves give the server a forward-looking view of a session.

4.3 Heterogeneity-aware KV-cache Arrangement

4.3.1 Adaptive Prompt re-layout

Prefix reuse stops at the first byte that differs from a cached prompt. The message a framework emits for a call site follows the declaration order of the function, which may place a changing value ahead of values that later calls could reuse. ANON’s adaptive re-layout permutes the components of each call message so that the bytes most likely to be




Figure 4. Example of adaptive re-layout, which increases prefix reuse by moving the more reusable components to the front, exposing shared prefix to the KV-cache.

resident already come first. The descriptor $\mathcal{T}_c$ (§4.2.1) and $\mathcal{A}$ (§4.2.2) determine a partial order over the components, and the server completes it into a layout with the token lengths it observes at runtime.

Transportable Components First, the dynamic bindings $\mathcal{B}_c$ can be permuted freely. The compiler lowers each argument into a self-contained component labelled by its parameter name, and the framework places these components in declaration order only because that is the order of the argument list; Re-layout may therefore choose any order of the bindings.

Second, part of the invariant specification $\mathcal{I}_c$ is movable as well. $\mathcal{I}_c$ is a simplified abstraction. More precisely, it consists system prompt, response schema, tool schema and instruction specification (Task description and parameter description generated by compiler). Among them, system prompt stays universally constant, and two schemas are digested by model separately. Instruction header is movable.

ANON therefore operates on the order of bindings and the placement of instruction specification. The following paragraphs describe how ANON chooses a layout within this space.

Layout decision algorithm Each call site sorts the bindings it carries by a program-wide key. Generally, *shared* bindings, which are carried by multiple call sites, are placed at the front. Then the bindings carried by a single call site.

Sorting private bindings. The bindings carried by a single call site are sorted in ascending lexicographic order of

$$K(b) = (\text{scoped}, \rho, \text{depth}, -\text{length}),$$

whose components are defined as follows: (i) *scoped*: the

lifetime of b ; 0 if it is session-wide ($S = \emptyset$) and 1 otherwise, so bindings that survive scope boundaries precede those reset by them. (ii) ρ : the stability rank defined in §4.2.1, from 1 for CONST to 5 for VOLATILE, so the more reusable class leads. (iii) *depth*: The number of transitions to the first call site that carries it from entrance, so values that enter a session earlier lead. (iv) *length*: the token length of the value of b , observed by the server; negated so longer bindings precede.

Ordering shared bindings. The ordering of shared bindings is derived from program-level structures. For every transition from q to q' , Static Analysis identifies how each shared binding of the two call sites evolves. We categorize these patterns into three groups and ordered in following sequence: (i) Both call sites carry it and its value does not change between the two calls. (ii) Both call sites carry it and the program only appends to its value. (iii) Program rewrites it between the two calls, or only one of the two call sites carries it.

ANON collects these requirements over all transitions and orders the shared bindings to satisfy them; when several bindings can go next, the one with the smallest (*-length, depth*) is placed first. Bindings with conflicting requirements are clustered and sorted by the same key.

Specification placement. The *instruction specification* is lowered from code by the compiler and is unique to each call site. Emitting it first would therefore prevent prefix sharing among different call sites, however many values they have in common. After sorting, ANON moves the specification behind the bindings whenever the leading binding is shared, so that sibling call sites can share large blocks of common bindings. A per-program switch keeps the specification in front when moving it is suspected of changing the model’s behavior.

Algorithm 1 summarizes the procedure.

4.3.2 Proactive KV-cache Scheduling

While adaptive prompt re-layout improves prefix reuse only by reshaping the request, ANON’s KV-cache scheduling proactively proactively arrange the placement of KV cache entries. Using the agent descriptor, the server decides for every served KV-cache sequence what to keep, what to give up, and what to load back ahead of time.

Identify Reusable KV-cache Since re-layout already places the stable and shared bindings at the front, a served prompt naturally divides into two segments: a reusable head and a volatile tail. In a full prompt, what behind the reusable head are all classified into volatile tail.

Reusable head consists two parts: a *static prefix* and a *session-shared extension*:

Algorithm 1 Layout decision algorithm

Require: Descriptor $(\{\mathcal{T}_c\}_{c \in C}, \mathcal{A})$, lengths ℓ , switch pin
Ensure: Layouts $\{L_c\}_{c \in C}$

```

1:  $J(b) \leftarrow (-\ell(b), \text{depth}(b))$ 
2:  $B_{sh} \leftarrow \{b : |\{c \in C : b \in \mathcal{B}_c\}| > 1\}$ 
3:  $R \leftarrow \emptyset$ 
4: for all  $e = (q, q') \in \text{EDGES}(\mathcal{A})$  do
5:    $B_e \leftarrow B_{sh} \cap (\mathcal{B}_q \cup \mathcal{B}_{q'})$ 
6:    $U_e \leftarrow \text{UNCHANGEDBINDINGS}(e, \mathcal{B}_q \cap \mathcal{B}_{q'})$ 
7:    $A_e \leftarrow \text{APPENDONLYBINDINGS}(e, \mathcal{B}_q \cap \mathcal{B}_{q'})$ 
8:   for all  $b \in B_e$  do
9:      $t_e(b) \leftarrow \begin{cases} 1 & b \in U_e \\ 2 & b \in A_e \\ 3 & \text{otherwise} \end{cases}$ 
10:  end for
11:   $R \leftarrow R \cup \{(b, b') \in B_e^2 : t_e(b) < t_e(b')\}$ 
12: end for
13:  $D \leftarrow \text{CLUSTERCONFLICTS}((B_{sh}, R))$ 
14:  $J_{\min}(S) \leftarrow \min_{b \in S} J(b)$ 
15:  $\sigma \leftarrow \text{TOPOSORT}(D, J_{\min})$  ▷ Min-key first
16:  $\pi_{sh} \leftarrow \langle \rangle$ 
17: for all  $S \in \sigma$  do
18:    $\pi_{sh} \leftarrow \pi_{sh} \parallel \text{SORT}(S, J)$ 
19: end for
20: for all  $c \in C$  do
21:    $\pi_c \leftarrow \langle b \in \pi_{sh} : b \in \mathcal{B}_c \rangle$ 
22:    $\pi_c \leftarrow \pi_c \parallel \text{SORT}(\mathcal{B}_c \setminus B_{sh}, K)$ 
23:    $s_c \leftarrow \text{INSTRUCTIONSPEC}(\mathcal{I}_c)$ 
24:   if  $\neg pin \wedge (\mathcal{B}_c \cap B_{sh} \neq \emptyset)$  then
25:      $L_c \leftarrow \pi_c \parallel \langle s_c \rangle$ 
26:   else
27:      $L_c \leftarrow \langle s_c \rangle \parallel \pi_c$ 
28:   end if
29: end for
30: return  $\{L_c\}_{c \in C}$ 

```

- **Static Prefix** is what the call site emits without runtime inputs, and its values outlives single session. It includes the invariant specifications $\mathcal{I}_c$ together with the CONST bindings that re-layout has placed right behind them.
- **Session-shared Extension** is the run of bindings behind the static prefix that a later call of the same session will emit again at the same position.

The two parts are derived from agent program itself. $\mathcal{H}_c$ names the provenance of each binding, and the automaton $\mathcal{A}$ names the transitions between call sites. Additionally with the decided layout of each call-site, we determine the reusable head by jointly analyze them.

Predict Future Reuse Because the workflow automaton $\mathcal{A}$ already exposes all potential next calls, our prediction focus on the following three aspects: (i) The probability of the arrival. (ii) When will the call arrive. (iii) Which prefix of its prompt is already known.

Probability of arrival. We estimate the probability of the next call by a loop-tagged table. As the compiler side emits the workflow graph, ANON adopts Bourdoncle-style (Bourdoncle, 1993) weak topological ordering to discover all loops within the agent, so that each call-site q is assigned the chain of loops enclosing it. The server counts the iterations of each loop per session, so every call carries a tuple ζ of the current iteration numbers of the loops enclosing its call-site, outermost first. We refine the transition and termination probabilities of $\mathcal{A}$ by this tuple and define the table as

$$\mathcal{P}(q' \mid q, \zeta) \quad \mathcal{F}(q, \zeta)$$

where ζ is the loop indices; $\mathcal{P}$ stands for the probability that a call-site q in context ζ is followed by a call-site q' ; $\mathcal{F}$ stands for the probability that the session ends at call-site q under context ζ . With an empty ζ they reduce to P and F of §4.2.2.

The table is a frequency count that the server updates across sessions. When a call arrives, its call-site is counted as the outcome of the previous call-site under the previous call's ζ , and termination is counted when the session closes. Each observation is also counted under every suffix of ζ down to the empty tuple, so that for the same q the table holds one row per level of context, from the fully specified iteration to the pooled count over all iterations. After a call at q lands, we compute the probability of each outcome o by Witten–Bell interpolation from the empty context ζ_0 to the full context $\zeta_m = \zeta$, where o is a successor of q or a session termination:

$$\hat{P}_k(o) = \lambda_k \frac{c_k(o)}{n_k} + (1 - \lambda_k) \hat{P}_{k-1}(o),$$

$$\hat{P}_0(o) = \frac{1}{|\text{succ}(q)|}, \quad \lambda_k = \frac{n_k}{n_k + d_k},$$

where $\text{succ}(q)$ is the set of static successors of q (including termination, which is treated as at an exit call-site). ζ_k is the suffix of ζ that keeps only the innermost k indices, and the three counts are read from the row (q, ζ_k) of the table: $c_k(o)$ is how many times outcome o followed a call at q in context ζ_k , $n_k = \sum_o c_k(o)$ is the number of observations in that row, and d_k is the number of distinct outcomes with $c_k(o) > 0$.

Then,

$$\mathcal{P}(q' \mid q, \zeta) = \hat{P}_m(q') \quad \mathcal{F}(q, \zeta) = \hat{P}_m(\text{end})$$

A level with no observation is skipped. The weight λ_k grows with the support of a level, so a well-observed iteration

dominates the prediction while a sparse one defers to the pooled behavior of q .

Estimate call arrival time. The time of arrival is conservatively approximated based on record statistics. The workflow automaton records $D_{q,q'}^{\text{gap}}$ per transition: each time a call at q' arrives, the server appends the interval between the reply of the preceding call at q and this arrival. The interval is the agent’s own time spent in its code between two calls.

We estimate the gap before the next call as:

$$\hat{g} = t_{50} - \frac{1}{2}(t_{90} - t_{50})$$

where t_{50} and t_{90} are the 50th and 90th percentiles of the gap samples on the transition (q, q') ; the predicted call is then expected $\hat{g}$ after the reply of the current call. We introduce the second term to make sure a transition with more sparse distribution of gap receives a tighter estimation.

Construct prompt ahead of time. As indicated in 4.3.2, the prompt of the known successor q' is reconstructible, or at least partial-reconstructible. ANON builds its prompt in the layout $L_{q'}$, using the values that the session already carries. The construction processes bindings in the order of $L_{q'}$ while consulting their evolution class, deducing value accordingly. The construction stops at the first binding with no known origin. Then the matched prefix is sent to kv-planner for speculative preparation.

Lifecycle aware KV-cache scheduling As we discussed in 3, different components have different lifetime, sharing scope and evolution patterns. Therefore, relying on recency based LRU would lead to suboptimal performance. ANON adopts a banded priority and steers KV-cache upon it.

Definition of Priority Bands. We categorize KV-cache entries into four bands, from first evicted to last:

1. **RETIRED:** the volatile tail of prompt belonging to completed sessions.
2. **TRANSIENT:** the volatile tail of a served prompt of a live session. The session that produced them will not present them again.
3. **UNCLASSIFIED:** entries that analysis cannot categorize, such as the prompts of a program that was never registered. They keep the engine’s default treatment.
4. **PLANNED:** prefixes that a predicted call will reuse. Inside this band entries are ordered by the expected value of keeping them resident, defined next.

Eviction empties the lowest band first and falls back to recency only within a band. Generally, KV-cache that will likely be reused left memory later.

Band Based KV-cache scheduling. As we mentioned above, eviction empties the lowest band first and falls back to recency only within a band, so KV-cache that will likely be reused left memory later. On the host tier, however, it keeps its recency order except that retired segments go first, so a head the device gave up under pressure remains reloadable while it may still be needed. Additionally, KV-cache entries marked as *TRANSIENT* won’t be admitted into host as they rival memory space with reusable KV-cache entries.

Bands are assigned at three moments of a session’s lifecycle. (1) **When a served prompt enters the cache**, ANON splits it at its reusable head (Section 4.3.2). The volatile tail and the reply generated after it become *TRANSIENT*; the head stays *UNCLASSIFIED* until the next prediction refresh finds a successor call that starts with it, which promotes it to *PLANNED*. (2) **When a later call arrives**, the transition it takes may rewrite bindings that earlier heads carried. Each affected head is cut back to the prefix that survives the rewrite, and the bytes past it are demoted to *TRANSIENT*. (3) **When the session ends**, the volatile tails of every prompt it served drop to *RETIRED*. Their heads stay *UNCLASSIFIED* and age out by recency, since other sessions of the same program may still share them.

These bands are mirrored directly into the radix tree’s eviction order. The band is the primary key of the eviction heap and recency breaks ties.

Prediction-based KV-cache protection Based on prediction, kv-planner of ANON conducts protection. The duty of it is to ensure required prefix resides on the device, or at least on host. Each prediction is summarized into a tuple of reconstructed prefix, arrival probability and estimated deadline:

$$\text{Prediction} = (\pi_{q'}, \mathcal{P}(q' | q, \zeta), t_{\text{reply}} + \hat{g})$$

where $\pi_{q'}$ is the prefix reconstructed for q' and t_{reply} is the time the current call replied. Predictions are refreshed at every call, and are dropped when the predicted calls arrive or expires. Protection acts on this tuple in two ways. *Retention* keeps $\pi_{q'}$ from leaving the device, and *prefetch* restores the part of $\pi_{q'}$ that has already been demoted to the host.

Retention. The planner ranks the predicted reusable KV-cache entries by the value of keeping $\pi_{q'}$ resident at every refresh:

$$\mathcal{D}(q') = \frac{\mathcal{P}(q' | q, \zeta) \cdot |\pi_{q'}|_{\text{cached}}}{\max(t_{\text{reply}} + \hat{g} - t_{\text{now}}, \tau_{\min})}$$

where $|\pi_{q'}|_{\text{cached}}$ is the number of tokens of $\pi_{q'}$ cached (both on host and device) and $\tau_{\min}$ is a floor that keeps a call due now from dominating. ANON’s kv-planner sort all predictions by $\mathcal{D}$ in descending order, and admits them into

PLANNED until the band covers a fixed fraction of the device tier. Within the band, their $\mathcal{D}$ add up on every radix cache node, making shared head outranks a private one.

Prefetch. When part of $\pi_{q'}$ has been demoted to the host, the planner prefetches it back to the device ahead of the estimated arrive time. A prefetch release deadline is therefore defined as:

$$t_{\text{release}} = t_{\text{reply}} + \hat{g} - \frac{|\pi_{q'}|_{\text{host}}}{B_{\text{h2d}}} - \frac{|\pi_{q'}|_{\text{uncached}}}{B_{\text{prefill}}}$$

where B_{h2d} and B_{prefill} are the measured host-to-device copy rate and prefill rate. Among released predictions the earliest deadline goes first, and break ties by expected saving:

$$\text{Saving} = \mathcal{P}(q' | q, \zeta) \cdot |\pi_{q'}|_{\text{cached}}$$

A prefetch takes only free device space after reclaiming *RETIRED* and *TRANSIENT* entries, so it never interferes with real requests; when no space is free, prefetching pauses briefly but still prevent predicted KV-cache entries from being completely evicted. A prefetch is released when the engine is idle or its deadline t_{release} approaches.

5 IMPLEMENTATION

We implemented ANON on top of Jac byLLM and SGLang. The serving layer of ANON sits upon model engine and ensure KV-Cache entries are scheduled with respect to their heterogeneity. The compiler side static analysis construct essential metadata from compiler-IR, sending the them to serving layer of ANON. Necessary modifications of SGLang HiCache engine for ANON’s KV-cache arrangement are also made based on demand.

6 EVALUATION

In this section, we evaluate ANON on multiple representative agent workloads, and ablate each components.

6.1 Environmental Settings

6.1.1 Testbed and Model

We conduct all experiments on one NVIDIA RTX 3090 GPU using Qwen3-14B-AWQ. We allocate 5.2 GB of GPU memory to the device KV-cache pool and 10 GB of CPU memory to the host KV-cache pool.

6.1.2 Baselines

We compare ANON against the following serving systems, together with two ablated variants of ANON. We also run each established serving system with ANON’s adaptive prompt re-layout enabled: (i) **Vanilla SGLang** (Zheng et al., 2024) uses SGLang’s HiCache with LRU eviction.

(ii) **CacheScout** (Zhang et al., 2026) identifies agents from their request prefixes and learns inter-agent execution transitions online. It uses the learned transition patterns to prioritize cached KV entries and prefetch the prefixes of predicted successor agents. (iii) **Continuum** (Li et al., 2025) retains each agent program’s KV cache for an adaptively selected time-to-live (TTL). The TTL is derived from observed inter-invocation delays and expected cache reuse, allowing the system to preserve reusable context while the agent executes tools. (iv) **KVFlow** (Pan et al., 2025) uses an Agent Step Graph to build a workflow-aware KV cache management framework for agentic workloads. (v) **ANON-KV-only** enables proactive KV-cache scheduling but disables adaptive prompt re-layout. (vi) **ANON-Relayout-only** enables adaptive prompt re-layout but replaces proactive KV-cache scheduling with LRU eviction.

6.1.3 Workloads

We evaluate ANON on four representative agent workloads, written with ByLLM (Dantanarayana et al., 2025) in Jac: a HoVer-style (Jiang et al., 2020) fact-checking agent, a HumanEval (Chen et al., 2021) coding agent, a BFCL (Patil et al., 2025) tool-use agent, and a FinanceBench (Islam et al., 2023) financial-analysis agent. We further evaluate a mixed workload comprising all four agent types to assess ANON’s end-to-end performance under concurrent, heterogeneous demand. The mixed workload runs 20 concurrent lanes, five per agent type; when a lane finishes a session, it immediately starts the next one.

6.2 End-to-End Performance of ANON

We tested ANON under different number concurrent workflows, and a mixed workloads. Across these workloads, we compare average job completion time(JCT), average time-to-first-token(TTFT) and throughput(session/min) for the mixed workload.

Performance on mixed workload As shown in table 1, ANON raises throughput to $1.44\times$ that of Vanilla SGLang and reduces average TTFT and JCT by 34% each. Continuum and CacheScout perform close to Vanilla SGLang on all three metrics. The gain of ANON comes from a higher cache hit rate in both device and host memory.

Average JCT Across the four workloads, ANON reduces average JCT by 6–48% compared with Vanilla SGLang (Figure 5), and by 34% on the mixed workload (Table 1). When concurrency is low, adaptive prompt re-layout accounts for nearly all of the gain. As concurrency grows, the gain of re-layout alone shrinks, and on coding and BFCL it turns into a slowdown; proactive KV-cache scheduling carries the gain: it evicts volatile tails first and protects the reusable heads that predicted calls will need. On fact check, scheduling




Figure 5. Average TTFT speedup (top) and average JCT speedup (bottom) over Vanilla SGLang under varying numbers of concurrent workflows.

Table 1. Performance on the mixed workload. TTFT and JCT are weighted average over three agents by their lane count. Throughput is calculated under steady load.

	Throughput	TTFT	JCT	Hit rate (%)		
System	(sess/min)	(s)	(s)	Device	Host	Total
<i>Original prompt order</i>						
Vanilla SGLang	2.81	17.5	555	11.3	1.3	12.6
CacheScout	2.81	17.1	543	13.1	0.3	13.4
Continuum	2.81	17.3	552	11.6	1.4	13.0
KVFlow	2.91	17.7	553	13.1	0.6	13.6
ANON-KV-only	3.07	16.3	514	14.1	7.2	21.3
<i>With adaptive prompt re-layout</i>						
ANON-Relayout-only	3.31	14.1	435	8.7	29.5	38.1
CacheScout	3.45	13.1	418	23.6	18.7	42.2
Continuum	3.38	14.0	437	8.7	27.9	36.6
KVFlow	3.45	13.2	412	20.8	23.1	43.9
ANON w/o prefetch	3.71	12.3	382	16.7	32.6	49.4
ANON	4.04	11.5	364	17.3	33.4	50.7

contributes from the lowest concurrency.

Except for an 8–13% reduction on BFCL at $c = 4$, Continuum and CacheScout perform close to Vanilla SGLang. Both exploit predictability between invocations but treat each prompt as an opaque token sequence. In the original prompt order, volatile bindings often precede reusable ones, making prefixes hold little that later calls can reuse. ANON-KV-only, which schedules on the original prompt order, shows the same limit.

Table 2. Task accuracy (%) of systems serving the original and the re-laid prompts, and p -values of a paired permutation test.

Workload	Original	Re-laid	Δ (pp)	p
Fact check	20.0	20.4	+0.4	0.95
Coding	92.1	91.8	−0.2	0.66
BFCL	7.2	6.4	−0.8	1.00
Finance	34.6	35.4	+0.8	0.82

Average TTFT Across the four workloads, ANON reduces average TTFT by 12–78% compared with Vanilla SGLang (Figure 5), and by 34% on the mixed workload (Table 1). As with JCT, adaptive prompt re-layout accounts for nearly all of the gain at low concurrency, and proactive KV-cache scheduling carries it as concurrency grows. On fact check and BFCL, and on finance at low concurrency, the TTFT reduction exceeds the JCT reduction, because a reused prefix skips its prefill, whereas JCT also includes decoding and tool execution which the KV-cache does not shorten. Continuum and CacheScout show the same limit as in JCT.

Task Quality We therefore compare task accuracy between the five systems that serve the original prompts and the that serve the re-laid prompts, using a paired permutation test over inputs (Table 2). On all four workloads the difference is not significant ($p \geq 0.66$) and no consistent direction of change. The performance gains of ANON therefore do not come at the cost of task quality.

6.3 Cache Efficiency of ANON

Figure 6 shows the cache hit rate of prompt tokens at the same concurrency levels as Figure 5. On coding and BFCL,

	Fact check					Coding					BFCL				Finance				
Vanilla	30	15	11	8	8	23	20	20	20	20	13	9	9	8	7	4	5	7	6
Continuum	27	15	12	8	8	22	21	20	21	22	14	9	9	8	5	4	5	5	3
Continuum+Relayout	46	32	29	22	19	56	54	33	16	13	69	64	37	14	64	48	11	5	4
CacheScout	31	16	12	10	9	22	22	21	21	21	17	10	9	8	6	4	5	5	3
CacheScout+Relayout	47	33	27	23	22	56	55	34	21	19	69	64	38	14	64	48	13	5	4
KVFlow	35	19	13	9	8	22	21	21	21	21	14	9	9	8	7	4	5	7	6
KVFlow+Relayout	54	40	31	25	22	57	55	38	27	23	69	64	46	20	64	50	15	5	4
KVOnly	36	29	20	14	12	27	27	27	25	25	14	9	8	8	4	7	6	7	6
Relayout	46	32	28	23	20	57	54	28	15	14	69	64	36	12	64	51	9	7	3
Ours	62	46	35	27	25	56	57	57	56	51	69	64	58	25	63	56	26	18	14
	4	6	8	10	12	24	32	40	48	52	4	8	16	24	4	8	12	16	20
	Concurrency																		

Figure 6. Cache hit rate (%) of prompt tokens under varying numbers of concurrent workflows.

adaptive prompt re-layout alone brings the hit rate at low concurrency to that of ANON: it places the stable and shared bindings at the front, where they form a reusable head that later calls hit. As concurrency grows, the hit rate of ANON-Relayout-only drops quickly, because LRU eviction does not distinguish reusable heads from volatile tails. With proactive KV-cache scheduling, ANON keeps a higher hit rate than ANON-Relayout-only as concurrency grows, and on coding its hit rate stays nearly flat. On fact check, ANON’s hit rate exceeds that of ANON-Relayout-only from the lowest concurrency. Continuum and CacheScout stay close to Vanilla SGLang at every concurrency. ANON-KV-only stays well below ANON, since under the original prompt order the reusable head it can protect is short. The mixed workload shows the same pattern (Table 1).

6.4 Ablation Studies

We ablate ANON on the mixed workload by enabling its components one at a time (Table 1). Besides ANON-KV-only and ANON-Relayout-only, ANON w/o prefetch keeps adaptive prompt re-layout and proactive KV-cache scheduling but disables prefetch. Figures 5 and 6 show ANON-KV-only and ANON-Relayout-only across concurrency levels.

Adaptive prompt re-layout alone reduces JCT by 22% compared with Vanilla SGLang and raises throughput to $1.18\times$. Adding proactive KV-cache scheduling without prefetch reduces JCT by a further 12% and raises throughput to $1.32\times$. Prefetch reduces JCT by another 5% and brings throughput to $1.44\times$.

The two mechanisms depend on each other. Applied to the original prompt order, proactive KV-cache scheduling reduces JCT by only 7% (ANON-KV-only); applied after re-layout, it reduces JCT by 16% relative to ANON-Relayout-only. The combined 34% reduction is larger than the sum of the two individual reductions. Re-layout places the stable

and shared bindings at the front, so a served prompt divides into a reusable head and a volatile tail, and scheduling gets a reusable head to protect.

7 RELATED WORK

KV-cache reuse PagedAttention (Kwon et al., 2023) and SGLang (Zheng et al., 2024) reuse KV along an exact shared prefix, and hierarchical caches (Gao et al., 2024; Qin et al., 2025; SGLang Contributors, 2026) extend its capacity beyond device memory. Prompt Cache (Gim et al., 2024) and CacheBlend (Yao et al., 2025) reuse KV beyond the prefix, by fixing module positions in a schema or by selectively recomputing tokens of concatenated chunks. ANON keeps exact prefix matching and instead changes the prompt so that reusable components form the prefix; it is orthogonal to position-independent reuse.

Existing agent serving systems KVFlow (Pan et al., 2025), Continuum (Li et al., 2025), CacheScout (Zhang et al., 2026), PBKV (Zheng et al., 2026), and Pythia (Yu et al., 2026) predict future execution from step graphs, observed tool-call delays, learned transitions, or workflow histories, and use these predictions to retain and prefetch KV between calls.

Existing Exploitation of prompt structure Agent frameworks (Dantnarayana et al., 2025; Khattab et al., 2024; Furgale et al., 2026; Dantnarayana et al., 2026) know the structure of the prompts they generate. Where they arrange it for reuse, they apply one hand-chosen convention to every call site, such as NOAA’s static, history, and volatile regions. Parrot (Lin et al., 2024) passes application structure to the server through Semantic Variables that developers declare, but serves prompts in the layout the application wrote. SGLang’s frontend (Zheng et al., 2024) reorders prompt statements to lengthen the constant prefix, but has GPT-4 choose the order, which its authors note is not yet reliable.

8 CONCLUSION

The inherent structure and heterogeneity of programmatically generated prompts create opportunities for optimization. ANON identifies and exploits these opportunities rather than treating each prompt as a monolithic unit, thereby unlocking proactive optimizations that improve performance and efficiency. Our evaluation demonstrates that ANON’s methodology successfully creates a more efficient LLM serving system.

REFERENCES

Bourdoncle, F. Efficient chaotic iteration strategies with widenings. In *Formal Methods in Programming and Their*

- Applications*, volume 735 of *Lecture Notes in Computer Science*, pp. 128–141. Springer, 1993. doi: 10.1007/BFb0039704.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. d. O., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., Ray, A., Puri, R., Krueger, G., Petrov, M., Khlaaf, H., Sastry, G., Mishkin, P., Chan, B., Gray, S., Ryder, N., Pavlov, M., Power, A., Kaiser, L., Bavarian, M., Winter, C., Tillet, P., Such, F. P., Cummings, D., Plappert, M., Chantzis, F., Barnes, E., Herbert-Voss, A., Guss, W. H., Nichol, A., Paino, A., Tezak, N., Tang, J., Babuschkin, I., Balaji, S., Jain, S., Saunders, W., Hesse, C., Carr, A. N., Leike, J., Achiam, J., Misra, V., Morikawa, E., Radford, A., Knight, M., Brundage, M., Murati, M., Mayer, K., Welinder, P., McGrew, B., Amodei, D., McCandlish, S., Sutskever, I., and Zaremba, W. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Dantnarayana, J. L., Kang, Y., Sivasothynathan, K., Clarke, C., Li, B., Kashmira, S., Flautner, K., Tang, L., and Mars, J. Mtp: A meaning-typed language abstraction for ai-integrated programming. *Proceedings of the ACM on Programming Languages*, 9(OOPSLA2):1176–1204, 2025. doi: 10.1145/3763092.
- Dantnarayana, J. L., Kashmira, S., Tang, L., and Mars, J. SIGIL: Compiling agent skills into typed harnesses, 2026. URL <https://arxiv.org/abs/2607.27309>.
- DSPy Contributors. ChatAdapter: DSPy api reference. Online documentation, 2026. URL <https://dspy.ai/api/adapters/ChatAdapter/>. Accessed September 14, 2026.
- Furgale, P., Klingler, S., Nolan, J., Staats, M., Di Lorenzo, G., Martinez Abad, E., Schüller, C., Dinu, R., Devoto, A., Berard, P., Kaplun, G., Sarafian, E., Roveri, R., Derczynski, L., and Silveira Cabral, R. NVIDIA-labs OO Agents: Native Python Object-Oriented Agents, 2026. URL <https://arxiv.org/abs/2607.20709>.
- Gao, B., He, Z., Sharma, P., Kang, Q., Jevdjic, D., Deng, J., Yang, X., Yu, Z., and Zuo, P. Cost-efficient large language model serving for multi-turn conversations with CachedAttention. In *2024 USENIX Annual Technical Conference (USENIX ATC 24)*, pp. 111–126, Santa Clara, CA, 2024. USENIX Association. URL <https://www.usenix.org/conference/atc24/presentation/gao-b-in-cost>.
- Gim, I., Chen, G., seob Lee, S., Sarda, N., Khandelwal, A., and Zhong, L. Prompt Cache: Modular attention reuse for low-latency inference. In *Proceedings of Machine Learning and Systems*, volume 6, pp. 325–338, 2024. URL https://proceedings.mlsys.org/paper_files/paper/2024/hash/a66caa1703fe34705a4368c3014c1966-Abstract-Conference.html.
- Islam, P., Kannappan, A., Kiela, D., Qian, R., Scherrer, N., and Vidgen, B. Financebench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*, 2023.
- Jiang, Y., Bordia, S., Zhong, Z., Dognin, C., Singh, M., and Bansal, M. HoVer: A dataset for many-hop fact extraction and claim verification. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pp. 3441–3460, 2020.
- Khattab, O., Singhvi, A., Maheshwari, P., Zhang, Z., Santhanam, K., Vardhamanan, S., Haq, S., Sharma, A., Joshi, T. T., Moazam, H., Miller, H., Zaharia, M., and Potts, C. DSPy: Compiling Declarative Language Model Calls into State-of-the-Art Pipelines. In *International Conference on Learning Representations*, 2024.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626, New York, NY, USA, 2023. Association for Computing Machinery. doi: 10.1145/3600006.3613165. URL <https://arxiv.org/abs/2309.06180>.
- Li, H., He, R., Mang, Q., Zhang, Q., Mao, H., Chen, X., Zhou, H., Zhang, H., Cheung, A., Gonzalez, J., and Stoica, I. Continuum: Efficient and robust multi-turn llm agent scheduling with kv cache time-to-live, 2025.
- Lin, C., Han, Z., Zhang, C., Yang, Y., Yang, F., Chen, C., and Qiu, L. Parrot: Efficient serving of LLM-based applications with semantic variable. In *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, pp. 929–945. USENIX Association, 2024.
- Mars, J., Kang, Y., Daynauth, R., Li, B., Mahendra, A., Flautner, K., and Tang, L. The jaseci programming paradigm and runtime stack: Building scale-out production applications easy and fast. *arXiv preprint arXiv:2305.09864*, 2023.
- Pan, Z., Patel, A., Hu, Z., Shen, Y., Guan, Y., Li, W.-L., Qin, L., Wang, Y., and Ding, Y. KVFlow: Efficient prefix caching for accelerating LLM-based multi-agent workflows, 2025. URL <https://arxiv.org/abs/2507.07400>.
- Patil, S. G., Mao, H., Yan, F., Ji, C. C.-J., Suresh, V., Stoica, I., and Gonzalez, J. E. The berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models. In *Proceedings of the 42nd*

International Conference on Machine Learning, volume 267 of *Proceedings of Machine Learning Research*, pp. 48371–48392. PMLR, 2025.

Qin, R., Li, Z., He, W., Cui, J., Ren, F., Zhang, M., Wu, Y., Zheng, W., and Xu, X. Mooncake: Trading more storage for less computation—a KVCache-centric architecture for serving LLM chatbot. In *23rd USENIX Conference on File and Storage Technologies (FAST 25)*, pp. 155–170, Santa Clara, CA, 2025. USENIX Association. URL <https://www.usenix.org/conference/fast25/presentation/qin>.

SGLang Contributors. HiCache system design and optimization. SGLang documentation, 2026. URL https://docs.sglang.io/docs/advanced_features/hicache_design. Accessed September 22, 2026.

Yao, J., Li, H., Liu, Y., Ray, S., Cheng, Y., Zhang, Q., Du, K., Lu, S., and Jiang, J. CacheBlend: Fast large language model serving for RAG with cached knowledge fusion. In *Proceedings of the Twentieth European Conference on Computer Systems*, New York, NY, USA, 2025. Association for Computing Machinery. doi: 10.1145/3689031.3696098. URL <https://www.microsft.com/en-us/research/uploads/prod/2024/09/eurosys25-final999.pdf>.

Yu, S., Shu, J., Ni, Y., Qian, K., Li, X., Wang, Y., Zhang, J., Xu, Z., Yang, S., Zhu, L., Zhai, E., Lu, Q., Xing, J., Lu, Y., Jin, X., Liu, X., and Xu, H. Pythia: Exploiting workflow predictability for efficient agent-native LLM serving, 2026. URL <https://arxiv.org/abs/2604.25899>.

Zhang, R., Kim, C., Feng, S., Du, K., Liu, Y., Zhong, Y., Ching, C.-W., Jiang, J., and Hu, L. Learning agent execution for kv-cache management in agentic serving, 2026.

Zheng, H., Fu, F., Wu, J., Yuan, B., Zhang, Y., Wang, H., Zhu, Y., Yan, X., and Jiang, J. Efficient serving for dynamic agent workflows with prediction-based KV-cache management, 2026. URL <https://arxiv.org/abs/2605.06472>.

Zheng, L., Yin, L., Xie, Z., Sun, C., Huang, J., Yu, C. H., Cao, S., Kozyrakis, C., Stoica, I., Gonzalez, J. E., Barrett, C., and Sheng, Y. Sglang: Efficient execution of structured language model programs. In *Advances in Neural Information Processing Systems*, volume 37, 2024.